



Reichman University

Efi Arazi School of Computer Science

M.Sc. Program in Machine Learning and Data Science

Final Project Report

Permuted Dogs versus Cats

Student 1 ID: 123456789 **Student 2** ID: 987654321

Advisors: Shuli Finley (Reichman University) — Alon Oring (Reichman University)

June 2, 2026

Abstract

Spatial structure is central to natural-image classification, but pretrained models may retain class-discriminative information even when that structure is disrupted. This report studies binary Dogs-versus-Cats classification under fixed tile-wise permutations, motivated by structured-permutation hardness questions but evaluated as an empirical computer-vision testbed. Images from the Kaggle Dogs vs Cats dataset are divided into deterministic grid permutations across four grid settings and three permutation presets. The experiments compare frozen pretrained ResNet-18, DeiT-Tiny, and MLP-Mixer backbones with trained binary classification heads; evaluate ResNet-18 ablations involving augmentation, mixed-input training, classifier-head capacity, and curriculum strategies; and compute model-agnostic permutation hardness metrics joined to validation accuracy. Across the deterministic matrix, models remain far above chance, while accuracy decreases as tile count and permutation disruption increase. DeiT-Tiny is strongest in the frozen-backbone comparison, the ResNet-18 MLP head gives the largest average ablation improvement among the tested strategies, and adjacency destruction is the strongest predictor of lower ResNet-18 best validation accuracy among the tested hardness metrics. These findings are specific to fixed deterministic permutations and frozen pretrained representations. They do not establish that the models reconstruct or invert the permutation; high accuracy may also reflect surviving local texture, object-part, background, or other class-discriminative cues.

1 Introduction

Image classification models are typically evaluated on data whose spatial structure is intact. This report asks how much binary Dogs-versus-Cats classification performance survives when that structure is deliberately disrupted by a fixed tile-wise permutation. Each input image is split into an $N \times N$ grid, the tiles are rearranged by a deterministic mapping, and the classifier must still distinguish dogs from cats on a held-out validation set transformed by the same target permutation.

The main research question is: *under a fixed supervised train/validation split, how do grid size, deterministic permutation preset, and pretrained visual representation affect validation accuracy after tile-wise image permutation?* The report also asks two secondary, testable questions. First, which ResNet-18 training or head modifications improve performance relative to the matched frozen-backbone baseline? Second, do deterministic, model-agnostic permutation hardness metrics correlate with the observed ResNet-18 best validation accuracy across the final experiment matrix?

The task is motivated by permuted-puzzle problems and cryptographic hardness [1], but the cryptographic connection is motivational rather than directly tested. The present image setting is a relaxed empirical analogue: it uses natural images rather than graphs of functions, applies coarse tile-wise rather than pixel-wise permutations, and evaluates supervised classification rather than cryptographic recovery or security. The results should therefore be interpreted as evidence about this computer-vision testbed, not as evidence about cryptographic hardness.

This report makes four modest empirical contributions:

- A deterministic Dogs-versus-Cats tile-wise permutation benchmark over four grid settings and three deterministic permutation presets.
- A frozen-backbone comparison of convolutional, attention-based, and MLP-mixing visual representations.
- A ResNet-18 ablation study of augmentation, mixed-input training, classifier-head capacity, and curriculum strategies.
- A set of model-agnostic permutation hardness metrics joined to validation accuracy.

The implementation is organized as a reproducible notebook-based workflow. Part 1 trains model baselines on unpermuted images and deterministic tile-wise permutations at 4-by-4, 7-by-7, and 10-by-10 resolution. Part 2 tests improvement strategies for the same permuted setting, including patch shuffle, regular image augmentation, mixed original/permuted training, a larger ResNet-18 classification head, and two curricula. Part 3 computes permutation hardness metrics and joins them to validation accuracy so that image-independent permutation properties can be compared against model performance.

2 Related Works

2.1 Spatial Structure and Convolutional Inductive Bias

Convolutional neural networks are built around locality, weight sharing, and translation-related structure. These assumptions are well matched to ordinary natural images, where nearby pixels and local texture patterns often carry semantic information. ResNet architectures are especially relevant because residual connections

improve the optimization of deep convolutional models [3]. Tile-wise permutation disrupts this setting: it preserves the content inside each tile but breaks many adjacency relations between neighboring image regions. ResNet-18 is therefore a useful baseline for testing how far frozen convolutional features remain useful when local spatial continuity is partially destroyed.

2.2 Patch-Based and Attention-Based Vision Models

DeiT-Tiny and MLP-Mixer provide complementary comparisons because both operate on image patches rather than relying only on convolutional local neighborhoods. DeiT uses self-attention to exchange information across patch tokens [5], while MLP-Mixer alternates token-mixing and channel-mixing MLPs [4]. These architectures are relevant to tile-wise permutation because a permuted image can still contain informative patches even when their global arrangement is corrupted. However, the comparison in this report should not be read as a pure test of architectural inductive bias: the models also differ in pretraining procedure, parameter count, patch/tokenization behavior, and feature geometry.

2.3 Jigsaw Tasks, Patch Shuffling, and Spatial Pretexts

The experiment also relates to jigsaw-style and patch-shuffling vision tasks, where spatial rearrangement is used either as a corruption or as a self-supervised pretext objective [add citation]. In those settings, a model may be trained to infer correct spatial order or to learn representations from corrupted views. The present work is different: the goal is supervised Dogs-versus-Cats classification un-

der a fixed deterministic tile-wise permutation, not explicit reconstruction of the original image or prediction of the permutation. This distinction matters because high classification accuracy does not imply that the model has recovered the original layout.

2.4 Robustness, Augmentation, and Curricula

The ResNet-18 ablations connect to broader regularization and robustness ideas in deep learning [2]. Patch shuffle and regular image augmentations introduce additional stochastic variation during training. Mixed original/permuted training tests whether clean-layout examples help the classifier head learn cues that transfer to the permuted validation distribution. Curriculum learning variants test whether staged exposure to increasing corruption probability or increasing permutation difficulty improves optimization [add citation]. In this report, these methods are evaluated as practical interventions for a frozen pretrained ResNet-18, not as general claims about robustness under all spatial corruptions.

2.5 Cryptographic and Permuted-Puzzle Motivation

The motivating proposal cites permuted puzzles as a source of cryptographic hardness, where fixed permutations can make structured objects difficult to recover or classify [1]. This report does not solve or evaluate a cryptographic task. Instead, it borrows the idea of structured permutation as a controlled way to disrupt spatial organization in images. The cryptographic motivation provides conceptual background, while the empirical claims are limited to supervised classification on the Dogs vs Cats benchmark under deter-

ministic tile-wise permutations.

3 Data Overview

The dataset is Kaggle’s Dogs vs Cats image dataset, with labels inferred from filenames such as `cat.123.jpg` and `dog.456.jpg`. The repository expects labeled training images under `data/dogs-vs-cats/train`; the unlabeled Kaggle test folder is not used for validation accuracy. The supervised split is stratified and class-balanced, and the task is evaluated with validation accuracy on held-out labeled images.

For the unpermuted setting, each image is processed as a normal dog or cat image and used as the benchmark for post-permutation performance. For the permuted setting, each image is divided into a square tile grid and the same deterministic tile-wise permutation is applied across images for a run. The final grid settings are one tile / none, 16 tiles, 49 tiles, and 100 tiles. The one-tile/no-permutation baseline is repeated under the easy, medium, and hard labels only so that all parts use the same 4×3 grid-permutation matrix. Aggregate summaries should account for the fact that these three records correspond to the same unpermuted condition rather than three distinct permutations.

4 Methodology

The experimental pipeline is organized into three notebooks. Part 1 establishes baselines by training on the unpermuted task and on fixed deterministic tile-wise permutations. Part 2 reuses the Part 1 ResNet-18 baseline and trains improvement ablations on the same task family. Part 3 computes permutation descriptors and correlates them with the validation accuracy observed for the corresponding deterministic records.

The primary Part 2 and Part 3 model is ResNet-18 [3]; Part 1 also compares DeiT-Tiny and MLP-Mixer. Experiments save raw run outputs, aggregated CSV summaries, and figures under `outputs/results` and `outputs/figures`. Aggregation reports mean final validation accuracy, mean best validation accuracy, standard deviation across the three deterministic permutation labels at each tile count, and the number of runs. Validation accuracy is the main evaluation metric because the held-out validation set is class-balanced. Best validation accuracy is useful for comparing training behavior across runs, but it can be optimistic as a final performance estimate because it selects the best epoch using the validation set.

4.1 Training and Validation Sets by Experiment

All experiments use the labeled Kaggle Dogs vs Cats training folder as the only source of supervised examples. Labels are inferred from filenames, and the samples are split once with a stratified class-balanced split using seed 42, validation fraction 0.2, and test fraction 0.0. For the full dataset this gives 20,000 training images and 5,000 validation images: 10,000 cats and 10,000 dogs for training, and 2,500 cats and 2,500 dogs for validation. The validation split is held out from optimization and is not modified by training-only augmentations such as patch shuffle, regular augmentation, original/permuted sampling, or curriculum-stage sampling.

The curriculum ablations in Part 2 use stage-specific training loaders, but they do not create separate held-out validation splits per stage. Across both curricula, validation remains the same 5,000-image held-out loader for the final target run. Non-curriculum Part 2 ablations train for 10 epochs, whereas both curriculum ablations train for 30 epochs.

Part	1	Training set	Validation set
run			
1 tile / none		20,000 unpermuted images repeated under easy, medium, and hard labels.	5,000 unpermuted held-out images.
4-by-4 permutations		Same 20,000 images with easy, medium, and hard deterministic 4-by-4 permutations.	Same 5,000 held-out images with the matching deterministic permutation.
7-by-7 permutations		Same 20,000 images with easy, medium, and hard deterministic 7-by-7 permutations.	Same 5,000 held-out images with the matching deterministic permutation.
10-by-10 permutations		Same 20,000 images with easy, medium, and hard deterministic 10-by-10 permutations.	Same 5,000 held-out images with the matching deterministic permutation.

Table 1: Part 1 training and validation sets. The same split is reused for ResNet-18, DeiT-Tiny, and MLP-Mixer Base, and each remote run trains for 10 epochs.

This design should be interpreted cautiously: any difference between curriculum and non-curriculum records is partly confounded by training duration, so the reported results cannot isolate curriculum scheduling from additional optimization time.

4.2 Deterministic Permutation Functions

The final experiments no longer sample random tile permutations for the main matrix. Instead, every part uses one shared determin-

istic generator with three parameter presets: easy, medium, and hard. This makes comparisons across Parts 1, 2, and 3 reproducible and interpretable, because a given difficulty label refers to the same construction wherever it appears.

The generator is controlled by relative parameters: the fraction of tiles to swap, the maximum swap distance as a fraction of grid width, row and column shift fractions, and whether to apply a global reversal before the shifts and swaps. This matters because fixed operation counts are not comparable across grid sizes: two swaps affect 25% of the tiles in a 4-by-4 grid but only 4% of the tiles in a 10-by-10 grid. The easy preset therefore uses a small tile-count-scaled swap fraction and short local movement; medium increases the swap fraction and adds moderate row and column shifts; hard uses the largest swap fraction, broader movement radius, stronger shifts, and a global reversal. For the one-tile baseline, all three labels map to the unpermuted image; the labels are still saved so the experiment matrix remains exactly 4 grid settings by 3 presets.

To make the deterministic permutation matrix fully reproducible, the report should list the exact parameter values for each easy, medium, and hard preset, including swap fraction, maximum swap distance, row shift, column shift, and reversal flag. If the values are defined in code rather than in the report text, the corresponding configuration path should be cited explicitly.

4.3 Architecture Comparison and Frozen Backbones

To compare architectural inductive biases beyond ResNet-18, the Part 1 pipeline also supports three pretrained feature extractors: ResNet-18, DeiT-Tiny [5], and MLP-Mixer

[4]. All three are used in the same transfer-learning regime: the ImageNet-pretrained backbone is frozen and only the final binary classification head is optimized. This setting tests the robustness of pretrained representations plus a small trained classification head; it does not test full representation learning under permutation, and results may differ under full fine-tuning. It also makes the relevant training-capacity comparison the number of learnable head parameters, not the total number of parameters stored in the network.

The three models differ in how they exchange information across the image. ResNet-18 uses convolutions, giving it a strong local-neighborhood bias. DeiT-Tiny represents the image as patches and uses self-attention to let patches interact. MLP-Mixer also operates on image patches, but replaces attention and convolutions with two MLP operations: token mixing, which combines information across patches, and channel mixing, which recombines feature dimensions inside each patch representation. The comparison is informative but not fully controlled: the models differ not only in inductive bias, but also in pretraining data and procedure, parameter count, patch/tokenization behavior, and feature geometry.

The MLP-Mixer choice required special care. The smaller `mixer_s16_224` architecture would have been closer in total parameter count, but the available `timm` configuration does not provide pretrained weights for that variant. Freezing an untrained MLP-Mixer backbone would train only a classifier on random fixed features, which would not be a meaningful comparison to pretrained ResNet-18 and DeiT-Tiny. We therefore use the larger pretrained `mixer_b16_224.goog_in21k_ft_in1k` model. Although this increases the number of frozen parameters and may increase forward-pass




Figure 1: Simplified MLP-Mixer flow. Token mixing exchanges information across patches, while channel mixing recombines features within each patch representation.

cost, the number of optimized parameters remains comparable to the other models: 1,538 learnable parameters for MLP-Mixer Base versus 1,026 for ResNet-18 and 386 for DeiT-Tiny.

4.4 Improvement Strategies

Part 2 evaluates training strategies against the regular Part 1 ResNet-18 baseline. The main distinction is whether a strategy adds stochastic training views, mixes clean and permuted inputs, changes the classifier head, or changes the order in which the model sees easy and hard inputs. In all cases, validation is performed on the held-out validation images with the target permutation and without training-only perturbations.

For the counts below, let B be the configured batch size and let $N_{\text{train}} = 20,000$ be the number of training images. The ablations do not append files to the dataset. Instead, they replace or resample the ordinary training view during the forward pass, so the stored increase is zero even when every processed example is augmented or curriculum-scheduled.

4.4.1 Regular Part 1 Baseline

Main idea. This is the reference point for Part 2. ResNet-18 is trained normally on the target permutation, with no extra augmentation, no curriculum, and the standard loss.

Reason. The baseline is necessary because every ablation must be judged against ordinary training on the same task. Without it, an improvement strategy might look useful simply because the target permutation is easier than another run.

Deeper dive. The baseline uses the same train and validation split as Part 1. For a permuted run, both training and validation images are transformed by the same fixed tile permutation for that run. It has one training stage and does not increase the dataset size. If x_i is an input image, P_π is the target tile permutation, and f_θ is the classifier, the baseline minimizes

$$\frac{1}{N_{\text{train}}} \sum_{i=1}^{N_{\text{train}}} \ell(f_\theta(P_\pi(x_i)), y_i).$$

The number of added augmented images is 0 per batch and 0 per epoch; the number of processed training views remains B per batch and 20,000 per epoch.

4.4.2 Patch Shuffle

Main idea. Patch shuffle randomly rearranges smaller patches during training. The model therefore sees extra local spatial disruptions in addition to the fixed tile permutation.

Reason. This is relevant because the main task already tests whether class information survives spatial reordering. If extra random patch movement helps, it suggests that robustness to local layout disruption is useful for permuted Dogs versus Cats classification.

Deeper dive. The target permutation is still the run-level task definition, but each training batch can also receive random patch shuffling. Let S_ρ denote a random within-image patch permutation sampled independently during training. The effective training view is

$$\tilde{x}_i = S_{\rho_i}(P_\pi(x_i)),$$

and the loss is computed on $\ell(f_\theta(\tilde{x}_i), y_i)$. In the implementation the shuffle probability is 1.0, so a batch of size B contains B shuffled views and an epoch contains 20,000 shuffled views. These views are not appended to the dataset; they replace the current batch examples. The validation loader does not use patch shuffle, so reported accuracy measures the fixed target validation permutation.

4.4.3 Regular Augmentations

Main idea. Regular augmentations apply standard image-level stochastic transformations during training while keeping the target validation permutation fixed.

Reason. This tests whether generic image augmentation helps the frozen ResNet-18 head learn more robust dog-versus-cat cues under tile permutation, or whether the extra variation is mismatched to the spatial corruption used at validation time.

Deeper dive. The training set remains the same 20,000 images. For each example, an image-level transform E is sampled during training:

$$\tilde{x}_i = E(P_\pi(x_i)).$$

The strategy has one stage and trains for 10 epochs. It stores no additional images; it produces one stochastic view per processed training example. Validation uses the deterministic target permutation without training-only augmentation.

4.4.4 Mixed Original/Permuted Training

Main idea. Mixed original/permuted training exposes the model to clean images and target-permuted images in the same run.

Reason. If clean images preserve class structure that is useful for the permuted task, mixing them into training may stabilize the classifier. If the clean examples pull the head away from the target validation distribution, the strategy may hurt.

Deeper dive. The ablation samples each training view as original with probability $p_{\text{original}} = 0.5$ and target-permuted otherwise:

$$\tilde{x}_i = \begin{cases} x_i, & z_i < p_{\text{original}}, \\ P_{\pi}(x_i), & z_i \geq p_{\text{original}}, \end{cases}$$

where $z_i \sim \text{Uniform}(0, 1)$. This gives an expected 10,000 clean views and 10,000 permuted views per epoch. The validation loader remains fully target-permuted, so the reported accuracy measures transfer from the mixed training distribution to the target permutation.

4.4.5 ResNet-18 MLP Head

Main idea. This ablation keeps the frozen ResNet-18 backbone but replaces the linear classifier with a small MLP classification head.

Reason. The frozen backbone may still extract useful ImageNet features from permuted tiles, but the mapping from those features to dog/cat labels may be less linearly separable after spatial disruption. A small MLP head tests whether additional trainable capacity at the classifier stage improves the target task.

Deeper dive. The data, target permutations, validation loader, and 10-epoch non-curriculum schedule match the regular Part 1

baseline. The only intentional change is the classifier head, so improvements can be interpreted as evidence that head capacity matters under the frozen-backbone setup.

4.4.6 Corruption-Probability Curriculum

Main idea. The model starts with mostly intact images and gradually sees the target permutation more often. The corruption probability increases across four stages.

Reason. Jumping immediately into a hard permutation may make optimization unstable. This curriculum tests whether gradual exposure lets the model first learn dog-versus-cat cues and then adapt those cues to more frequent spatial corruption.

Deeper dive. Each stage uses the same 20,000 training images and target tile permutation, but changes the probability that the permutation is applied. For stage probability q_s , the sampled training view is

$$\tilde{x}_i = \begin{cases} P_{\pi}(x_i), & z_i < q_s, \\ x_i, & z_i \geq q_s, \end{cases}$$

$$z_i \sim \text{Uniform}(0, 1).$$

The stages use $q_s \in \{0.1, 0.3, 0.5, 0.8\}$, so the expected number of permuted views per stage epoch is $q_s N_{\text{train}}$: 2,000, 6,000, 10,000, and 16,000 respectively. In a batch of size B , the expected number of permuted views is $q_s B$. This does not store a larger dataset; it changes the sampled training view. There is no separate validation set for each stage; validation always uses the fixed target validation loader. Curriculum runs train for 30 epochs in the final saved results.

4.4.7 Permutation-Difficulty Curriculum

Main idea. The model starts on easier spatial structure and moves toward the fi-

nal harder permutation. Instead of changing probability, it changes the tile grid difficulty across stages.

Reason. Smaller or easier permutations preserve more of the original image structure. This experiment tests whether a model can learn useful class cues under mild disruption before facing the final, harder tile rearrangement.

Deeper dive. Each stage uses the same 20,000 training images but changes the tile grid used for training. If g_s is the stage grid size and π_s is the corresponding sampled or target permutation, then

$$\tilde{x}_i^{(s)} = \begin{cases} x_i, & g_s = 1, \\ P_{\pi_s}^{(g_s)}(x_i), & g_s > 1. \end{cases}$$

The schedule begins with original images and then progresses through 2-by-2, 3 tiles per side, and 4-by-4 permutations; for 7-by-7 and 10-by-10 target runs, it adds a final target-grid stage. Stage permutations use the same deterministic difficulty label as the target record. Each stage processes B scheduled-difficulty views per batch and 20,000 per stage epoch. The stored dataset size does not increase, validation remains the held-out loader for the final target run, and curriculum runs train for 30 epochs in the final saved results.

4.5 Tile Permutation Hardness Metrics

We quantify the difficulty induced by a tile permutation using three complementary, model-agnostic axes and a fourth combined score. Let the image be divided into an $N \times N$ grid. For each source tile i , let (r_i, c_i) denote its original row and column, and let (r'_i, c'_i) denote its destination after applying the permutation.

4.5.1 Adjacency Destruction Hardness

Main idea. This metric asks how much of the original neighborhood structure survives after the tile permutation. If tiles that were next to each other are no longer neighbors, the image becomes harder for a convolutional model that relies on local structure.

For every original rightward and downward neighboring relation, we test whether the same two source tiles remain adjacent in the same direction after permutation. The number of such canonical adjacency relations is

$$M = 2N(N - 1).$$

If $A_{\text{preserved}}$ is the number of preserved directed adjacencies, adjacency destruction is

$$A = 1 - \frac{A_{\text{preserved}}}{M}.$$

A low value means that much of the original local structure is preserved. A high value means that neighboring tile relations were mostly destroyed.

4.5.2 Global Tile Displacement

Main idea. This metric measures how far tiles move from their original positions. A permutation that only nudges tiles nearby should preserve more global shape than one that sends tiles across the image.

For each tile, we compute its Manhattan displacement in grid coordinates:

$$\delta_i = |r'_i - r_i| + |c'_i - c_i|.$$

The normalized global displacement is

$$D = \frac{\sum_i \delta_i}{D_{\text{max}}},$$

where D_{max} is the maximum possible total Manhattan displacement for an $N \times N$ grid. A low value indicates that tiles remain close to their original locations. A high value indicates large-scale spatial relocation.

4.5.3 Spatial Permutation Entropy

Main idea. This metric asks whether tile movement follows a simple pattern or many different movement patterns. More varied movement directions and distances mean the permutation is less orderly.

For each tile, we compute

$$\Delta_i = (r'_i - r_i, c'_i - c_i)$$

and assign it to a movement bin

$$b_i = (\text{direction}(\Delta_i), \text{distanceTier}(\Delta_i)).$$

The direction is one of

$$\{\text{stationary}, N, S, E, W, NE, NW, SE, SW\}.$$

For non-stationary tiles, the normalized Manhattan distance is

$$d_i = \frac{|r'_i - r_i| + |c'_i - c_i|}{2(N - 1)}.$$

The distance tier is near when $0 < d_i \leq 1/3$, medium when $1/3 < d_i \leq 2/3$, and far when $2/3 < d_i \leq 1$. Stationary tiles use a single stationary bin. If p_b is the empirical probability of movement bin b , the normalized entropy is

$$E = \frac{-\sum_b p_b \log p_b}{\log(N^2)}.$$

A low value means that movement is structured or concentrated in a small number of movement patterns. A high value means that the permutation contains diverse movement directions and distances, making it more disorderly.

4.5.4 Combined Hardness Score

Main idea. The combined score summarizes the previous metrics into one number so permutations can be ranked by an overall notion of difficulty.

The final score is a convex combination of the three normalized components:

$$C = w_a A + w_d D + w_e E,$$

with default weights

$$w_a = 0.5, \quad w_e = 0.3, \quad w_d = 0.2.$$

A low combined score represents an easier permutation with preserved structure, small movement, and low movement disorder. A high combined score represents a harder permutation across these complementary axes.

4.6 Reproducibility Checklist

The current report provides the main split sizes, validation protocol, model families, epoch counts for the final ablation matrix, and output directories. For full reproducibility, the report should also specify or link to the implementation values below. Unknown values should be filled from the code or notebooks rather than inferred.

5 Evaluation Protocol

Validation accuracy is the main metric throughout the report because the held-out validation set is class-balanced, with equal numbers of cats and dogs. Final validation accuracy reports the value at the end of the configured training run. Best validation accuracy reports the maximum validation accuracy reached during training and is useful for comparing training behavior, but it can be optimistic as a final performance estimate because it selects the best epoch on the validation set.

For Part 3, the hardness-correlation analysis is based on the 12 deterministic records in the final ResNet-18 matrix. The correlations should therefore be treated as exploratory descriptors of this controlled matrix rather than




Figure 2: Part 1 mean best validation accuracy as a function of tile count for the three frozen pretrained backbones. Higher tile counts correspond to finer grids and more opportunities for tile-wise spatial disruption.

definitive statistical estimates. No p-values, confidence intervals, repeated-seed variability estimates, or independent test-set measurements are reported.

6 Results

Table 9 summarizes the frozen-backbone architecture comparison by tile count, and Figure 2 visualizes the same trend across grid sizes.

Part 1 shows that all three frozen-backbone models remain far above chance under deterministic tile-wise permutations, but performance decreases as the grid becomes more disruptive. DeiT-Tiny is strongest under this frozen-pretrained setting, with mean best validation accuracy of 92.97% even at 100 tiles. MLP-Mixer Base is usually second, while ResNet-18 degrades most sharply as the easy, medium, and hard records are averaged at larger grids. The 100-tile mean best accuracy falls to 87.87% for ResNet-18, compared with 89.42% for MLP-Mixer Base and 92.97% for DeiT-Tiny.

Table 10 reports the ResNet-18 ablation summary over the 12 deterministic records, and




Figure 3: Part 2 best-validation-accuracy delta from the matched ResNet-18 baseline. Positive values indicate improvement relative to the same grid size and deterministic permutation preset.

Figure 3 shows each strategy’s delta from its matched baseline.

Part 2 gives a more selective picture than the earlier exploratory ablations. The ResNet-18 MLP head is the strongest improvement, adding 0.86 percentage points of mean best validation accuracy over the matched baseline. Corruption-probability curriculum is the next-best strategy at +0.39 percentage points. Permutation-difficulty curriculum is approximately neutral on average, while mixed original/permutated training and regular augmentations are slightly below baseline. Patch shuffle is clearly harmful in this final matrix, losing 2.40 percentage points on mean best validation accuracy.

Table 11 reports exploratory correlations between deterministic hardness metrics and ResNet-18 best validation accuracy. Figures 4–6 visualize the metric values and their relationship to grid size, difficulty label, and accuracy.

Part 3 provides a ranking lens for the deterministic permutations. All four hardness metrics are negatively correlated with ResNet-18 best validation accuracy over the 12 joined records. Adjacency destruction

is the strongest single metric, with Pearson correlation -0.9044 and Spearman correlation -0.9736, suggesting that broken local tile-neighborhood structure is especially predictive of lower validation accuracy within this matrix. The combined hardness score is also strongly negative, but it does not exceed adjacency destruction on this run matrix. Because only 12 deterministic records are joined, these correlations should be treated as exploratory rather than definitive.

7 Discussion

The proposal asks whether a neural network can classify post-permutation images and whether performance depends on tile resolution and permutation choice. The final deterministic matrix answers the first question affirmatively for this setting: every model and every ResNet-18 ablation remains well above chance, including the hard 10-by-10 cases. At the same time, the results show that the task is not trivial. Accuracy drops with stronger deterministic disruption, and the drop is largest for the frozen convolutional ResNet-18 baseline.

The architecture comparison suggests that patch-based pretrained representations remain effective under this specific corruption. DeiT-Tiny has the best aggregate performance, and MLP-Mixer Base also exceeds ResNet-18 at larger tile counts. This does not show that DeiT-Tiny is generally superior for permuted images, nor that attention or token mixing solves permutation hardness. The backbones are frozen, pretrained on natural images, and differ in several ways beyond architecture. A cautious interpretation is that class-discriminative information survives the deterministic tile-wise permutations and can be accessed by pretrained patch-based features plus a small classification head.

The Part 2 ablations show that simply adding spatial noise is not enough. Patch shuffle hurts the final matrix, possibly because the validation task uses one fixed deterministic permutation rather than arbitrary extra shuffling. The MLP head improves most consistently, suggesting that the frozen ResNet-18 features may contain useful information that is not optimally separated by a linear head. The corruption-probability curriculum gives a smaller positive gain, while the permutation-difficulty curriculum is nearly neutral; however, curriculum runs also use a longer training duration, so their gains cannot be attributed solely to curriculum design.

The hardness metrics address the proposal's third objective: ranking permutations with the same tile count. The negative correlations suggest that geometry-only permutation descriptors capture meaningful variation in ResNet-18 performance within the deterministic matrix. Adjacency destruction is particularly informative, matching the intuition that destroying local tile neighborhoods is harmful for a convolutional model. Because the analysis uses only 12 deterministic records, it should be viewed as a hypothesis-generating result for future experiments with more permutation samples.

Finally, high accuracy should not be interpreted as evidence that the models learned to undo the permutation. The experiments do not directly distinguish spatial reconstruction, compensation for a fixed mapping, and exploitation of surviving class-discriminative cues. Dogs-versus-Cats classification may remain solvable from local texture, object parts, background statistics, or other cues that persist even when global shape is disrupted.

8 Limitations

Several limitations constrain the strength and scope of the conclusions. First, no repeated training seeds are reported, so run-to-run variation is not quantified. Second, no independent labeled test set is used; the Kaggle unlabeled test folder is not used for validation accuracy. Third, the architecture comparison is confounded by pretraining, parameter counts, patch sizes, tokenization behavior, and model-specific feature representations. Fourth, the frozen-backbone setup limits conclusions about how fully trainable models would adapt to tile-wise permutations.

The experimental matrix is also limited. Curriculum results are partly confounded by longer training duration, since curriculum ablations train for 30 epochs while non-curriculum ablations train for 10 epochs. The deterministic matrix contains only three permutation presets per grid size, and the hardness-correlation analysis uses only 12 records. Finally, Dogs-versus-Cats may remain solvable from local texture, object-part, or background cues even when global structure is disrupted. The report therefore does not directly distinguish learning or compensating for a fixed permutation from exploiting remaining class-discriminative cues.

9 Conclusion and Future Work

This project implements a reproducible Dogs-versus-Cats testbed for studying fixed tile-wise image permutations. The concrete empirical findings are that fixed permutations reduce performance but do not eliminate class-discriminative signal; frozen pretrained DeiT-Tiny is most robust among the compared backbones in the final matrix; the

ResNet-18 MLP head is the strongest tested ablation; and adjacency destruction is the most predictive hardness metric among the tested metrics.

Within the scope of the experiments, these results imply that pretrained visual representations can preserve useful dog/cat information even when global image layout is disrupted by deterministic tile-wise permutations. They also suggest that local adjacency destruction is a useful model-agnostic descriptor of permutation difficulty for frozen ResNet-18. The results do not prove that models reconstruct, invert, or explicitly learn the permutation, and they do not establish general robustness to unseen permutations or cryptographic hardness.

Future work should include repeated training seeds, independent test evaluation, full fine-tuning, more permutation samples per grid size, transfer tests between permutation functions, and experiments on datasets where global shape is more necessary for classification. Overall, the study provides a reproducible empirical framework for analyzing classification under deterministic spatial permutations, while showing that stronger statistical validation and transfer experiments are needed before drawing broader conclusions about permutation robustness.

Part 2 strategy	Training set	Added training exposure	Validation set
Regular baseline	Reused Part 1 ResNet-18 runs.	None beyond the fixed target permutation.	Reused Part 1 held-out validation accuracy.
Patch shuffle	Same 20,000 training images and target permutation.	No appended files; each batch of size B produces B shuffled training views, for 20,000 shuffled views per epoch.	Same 5,000 held-out images with the target permutation only.
Regular augmentations	Same 20,000 training images and target permutation.	Standard stochastic training transforms produce one augmented view per example per epoch.	Same 5,000 held-out images with the target permutation only.
Mixed original/permuted	Same 20,000 training images.	Each training view is sampled as original or target-permuted, using $p_{\text{original}} = 0.5$.	Same 5,000 held-out images with the target permutation only.
ResNet-18 MLP head	Same 20,000 training images and target permutation.	No new images; the linear classifier is replaced by a small MLP head.	Same 5,000 held-out images with the target permutation.
Corruption-probability curriculum	Same 20,000 training images at every stage.	No appended files; stage probabilities 0.1, 0.3, 0.5, and 0.8 give expected permuted counts of 2,000, 6,000, 10,000, and 16,000 views per stage epoch.	One fixed target validation loader; no separate validation set per stage.
Permutation-difficulty curriculum	Same 20,000 training images at every stage.	No appended files; each stage processes 20,000 views per epoch at its scheduled difficulty: original, 2-by-2, 3 tiles per side, 4-by-4, and for 7-by-7 or 10-by-10 targets an additional final target-grid stage.	One fixed final-target validation loader; no separate validation set per stage.

Table 2: Part 2 training and validation sets by ablation. The final output matrix uses ResNet-18 with 4 grid settings crossed with easy, medium, and hard deterministic permutation records, for 12 records per ablation.

Part 2 strategy	Stages	Experiment sequence	Dataset increase / per-epoch views
Regular baseline	Part 1 1	Standard training on the target permutation.	No increase: 20,000 training examples per epoch.
Patch shuffle	1	Standard training, with random patch shuffling applied inside training batches.	0 stored increase; B shuffled views per batch and 20,000 per epoch.
Regular augmentations	1	Standard training, with generic image augmentations in the training transform.	0 stored increase; one stochastic view per example per epoch.
Mixed original/permuted	1	Standard training, sampling clean and target-permuted views with equal probability.	0 stored increase; expected 10,000 clean and 10,000 permuted views per epoch.
ResNet-18 head	MLP 1	Standard target-permutation training, replacing the linear head with an MLP head.	No increase: the examples are unchanged, only the trainable head changes.
Corruption-probability curriculum	4	Permutation probability stages: 0.1, 0.3, 0.5, 0.8.	0 stored increase; expected permuted views per epoch are 2,000, 6,000, 10,000, and 16,000 by stage.
Permutation-difficulty curriculum	1, 4, or 5	Baseline target: original only. Tiled targets: original, 2-by-2, 3 tiles per side, 4-by-4, and the final target grid when larger than 4-by-4, using the same easy/medium/hard function as the target record.	0 stored increase; 20,000 scheduled-difficulty views per stage epoch.

Table 3: Part 2 ablation experiment design. Non-curriculum ablations train for 10 epochs; both curriculum ablations train for 30 epochs. None of the ablations duplicate files on disk.



Figure 4: Part 3 example permutations and model-agnostic hardness metrics for the deterministic grid and difficulty settings. The figure is intended to make the easy, medium, and hard presets visually comparable across grid sizes.

Part 3 analysis	Training set	Validation source
Hardness metric computation	Not applicable; Part 3 does not train a model.	Reuses Part 1 ResNet-18 validation accuracies.
Metric-accuracy join	Not applicable; only saved CSV outputs are joined.	Matches each metric row to the corresponding Part 1 permutation ID.
Correlation analysis	Not applicable; no parameters are updated.	Computes correlations against the joined held-out validation accuracies.

Table 4: Part 3 data usage. Part 3 analyzes saved permutations and validation accuracies rather than creating new train or validation splits.

Model	Main operation	Inductive bias
ResNet-18	Convolution	Local texture and translation structure
DeiT-Tiny	Self-attention	Patch-to-patch interactions
MLP-Mixer Base	MLP mixing	Patch and feature mixing without attention

Table 5: Architectural comparison of the three frozen feature extractors.

Model	Learnable	Frozen	Total
DeiT-Tiny	386	5,524,416	5,524,802
ResNet-18	1,026	11,176,512	11,177,538
MLP-Mixer Base	1,538	59,111,472	59,113,010

Table 6: Parameter counts for the frozen-backbone comparison with two output classes. “Learnable” counts only the classifier head optimized during training; “Frozen” counts pretrained parameters used only in the forward pass.




Figure 5: Combined hardness score by grid size and deterministic difficulty label. The score aggregates normalized adjacency destruction, global displacement, and spatial entropy.




Figure 6: Adjacency destruction, grid size, and ResNet-18 best validation accuracy across the 12 joined deterministic records. Lower accuracy is associated with higher adjacency destruction in this exploratory matrix.

Metric	Dimension	What it captures
Adjacency destruction hardness	Local topology	Are original neighbors still neighbors?
Global tile displacement	Global geometry	How far did tiles move from their original positions?
Spatial permutation entropy	Permutation disorder	How diverse or irregular are tile movement patterns?
Combined hardness score	Aggregated hardness	Weighted summary of the three normalized axes.

Table 7: Tile permutation hardness metrics used in the analysis.

Item	Value to report
Dataset source and folder structure	Kaggle Dogs vs Cats; labeled images under <code>data/dogs-vs-cats/train</code> ; exact download/preparation steps [specify from implementation].
Train/validation split procedure	Stratified class-balanced split; validation fraction 0.2; test fraction 0.0; split seed 42.
Image preprocessing	Image resizing and cropping policy [specify from implementation]; normalization [specify from implementation].
Tile transform details	Tile boundary handling [specify from implementation]; interpolation method [specify from implementation]; exact easy/medium/hard preset parameters [specify from implementation].
Optimization	Batch size [specify from implementation]; optimizer [specify from implementation]; learning rate [specify from implementation]; weight decay [specify from implementation]; scheduler [specify from implementation].
Training duration	Remote Part 1 runs train for 10 epochs; non-curriculum Part 2 ablations train for 10 epochs; curriculum Part 2 ablations train for 30 epochs.
Model configuration	Model checkpoint names; frozen versus trainable parameter policy; classification-head definitions; MLP-head architecture [specify from implementation where not already stated].
Augmentation and curricula	Patch-shuffle definition; regular augmentation definitions; mixed-input probability; curriculum stage definitions and stage epoch allocation [specify from implementation where not already stated].
Compute environment	Hardware, software versions, and deterministic settings [specify from implementation].
Artifacts	Code repository or notebook paths; output CSV locations; output figure locations; checkpoint paths [specify from implementation].

Table 8: Reproducibility checklist. Bracketed fields mark implementation details that should be reported from the code or notebooks rather than reconstructed from memory.

Model	1 tile best	16 tiles best	49 tiles best	100 tiles best
ResNet-18	97.89%	94.25%	91.69%	87.87%
DeiT-Tiny	98.75%	97.23%	97.08%	92.97%
MLP-Mixer Base	98.52%	95.97%	95.00%	89.42%

Table 9: Part 1 mean best validation accuracy by model and tile count, aggregated across the easy, medium, and hard deterministic records at each grid size.

Part 2 setting	Mean final acc.	Mean best acc.	Mean best delta vs. baseline
Regular Part 1 baseline	92.85%	92.92%	0.00 pp
Patch shuffle	89.82%	90.52%	-2.40 pp
Regular augmentations	92.36%	92.58%	-0.35 pp
Mixed original/permuted	92.76%	92.83%	-0.10 pp
ResNet-18 MLP head	93.41%	93.79%	+0.86 pp
Corruption-probability curriculum	93.08%	93.31%	+0.39 pp
Permutation-difficulty curriculum	92.87%	92.98%	+0.05 pp

Table 10: Part 2 ResNet-18 ablation summary over the 12 deterministic records. The delta column compares each record’s best validation accuracy against the matched regular Part 1 ResNet-18 baseline and then averages those deltas.

Hardness metric	Pearson	Spearman
Global displacement	-0.7607	-0.8169
Adjacency destruction	-0.9044	-0.9736
Spatial entropy	-0.6942	-0.8310
Combined hardness	-0.8477	-0.9296

Table 11: Part 3 correlations between ResNet-18 best validation accuracy and permutation metrics over the 12 deterministic records.

References

- [1] E. Boyle, J. Holmgren, and M. Weiss. Permuted puzzles and cryptographic hardness. In *Theory of Cryptography*, pages 465–493, 2019.
- [2] I. Goodfellow, Y. Bengio, and A. Courville. *Deep Learning*. MIT Press, 2016.
- [3] K. He, X. Zhang, S. Ren, and J. Sun. Deep residual learning for image recognition. In *IEEE Conference on Computer Vision and Pattern Recognition*, pages 770–778, 2016.
- [4] I. O. Tolstikhin, N. Houlsby, A. Kolesnikov, L. Beyer, X. Zhai, T. Unterthiner, J. Yung, A. Steiner, D. Keysers, J. Uszkoreit, M. Lucic, and A. Dosovitskiy. MLP-Mixer: An all-MLP architecture for vision. In *Advances in Neural Information Processing Systems*, 2021.
- [5] H. Touvron, M. Cord, M. Douze, F. Massa, A. Sablayrolles, and H. Jégou. Training data-efficient image transformers and distillation through attention. In *International Conference on Machine Learning*, pages 10347–10357, 2021.