



Reichman University

Efi Arazi School of Computer Science

M.Sc. Program in Machine Learning and Data Science

Final Project Report

Permuted Dogs versus Cats

Student 1 ID: 123456789 **Student 2** ID: 987654321

Advisors: Shuli Finley (Reichman University) — Alon Oring (Reichman University)

May 25, 2026

Abstract

This project studies binary Dogs versus Cats classification after fixed tile-wise permutations are applied to each image. The proposal frames the task as a relaxed image-domain analogue of permuted-puzzle problems with cryptographic motivation: instead of pixel-wise permutations over function graphs, images are divided into coarse grids and their tiles are rearranged before classification. We implement a notebook-first experimental pipeline around the Kaggle Dogs vs Cats dataset, train ResNet-18 baselines on unpermuted and permuted images, evaluate augmentation and curriculum strategies, and define model-agnostic permutation hardness metrics. In the current experiments, the unpermuted ResNet-18 baseline reaches 86.5% validation accuracy, while 3-by-3 tile permutation performance averages 83.7%. The best final 3-by-3 ablation is permutation-difficulty curriculum training at 89.4%, and hardness metrics show negative Pearson correlations with validation accuracy for adjacency destruction, spatial entropy, and combined hardness. These results suggest that tile permutations preserve enough class signal for convolutional networks to classify cats and dogs, while permutation structure still measurably affects difficulty.

1 Introduction

Image classification models are typically evaluated on data whose spatial structure is intact. This project asks what happens when that structure is deliberately disrupted before learning. Each input image is split into an $N \times N$ grid, the tiles are permuted by a fixed mapping, and the classifier must still distinguish dogs from cats. The task is motivated by a proposal for an intermediate experiment connected to permuted puzzles and cryptographic hardness [1]. The image setting relaxes the original cryptographic problem in two ways: it uses natural images rather than graphs of functions, and it applies tile-wise permutations rather than pixel-wise permutations.

The project has three main objectives. First, it measures how tile resolution affects the feasibility of classification after permutation. Second, it asks whether a neural network succeeds by learning or compensating for the permutation, or whether enough class-discriminative signal remains directly visible in the permuted images. Third, it develops ways to rank permutations according to their induced difficulty for a fixed dataset and tile count.

The implementation follows a reproducible, notebook-first workflow. Part 1 trains a ResNet-18 baseline on unpermuted images and 3-by-3 tile permutations. Part 2 tests improvement strategies for the permuted setting, including image augmentations and curricula. Part 3 computes permutation hardness metrics and joins them to validation accuracy so that image-independent permutation properties can be compared against model performance.

2 Related Works

The motivating proposal cites permuted puzzles as a source of cryptographic hardness, where fixed permutations can make structured objects difficult to recover or classify [1]. This project does not attempt to solve the cryptographic problem directly. Instead, it uses the Dogs vs Cats setting as an empirical stepping stone for studying how much semantic information survives structured spatial disruption.

Convolutional neural networks are a natural starting point because they exploit locality and translational structure in images. ResNet architectures are especially relevant because residual connections improve optimization of deep convolutional models [3]. Since tile permutations partially destroy local adjacency while preserving tile content, ResNet-18 provides a useful baseline for testing whether standard convolutional inductive biases remain helpful.

The ablations also connect to standard regularization and robustness ideas in deep learning [2]. Random erasing, patch shuffling, same-label CutMix, and combined augmentations are used to expose the model to local occlusion or spatial disruption. Curriculum learning variants gradually control corruption probability or permutation difficulty, testing whether staged exposure to harder permuted inputs improves validation accuracy.

3 Data Overview

The dataset is Kaggle’s Dogs vs Cats image dataset, with labels inferred from filenames such as `cat.123.jpg` and `dog.456.jpg`. The repository expects labeled training images under `data/dogs-vs-cats/train`; the unlabeled Kaggle test folder is not used for validation accuracy. The task is binary classifi-

cation, and the saved experiment outputs report validation accuracy after the configured training runs.

For the unpermuted setting, each image is processed as a normal dog or cat image and used as the benchmark for post-permutation performance. For the permuted setting, each image is divided into a square tile grid and the same tile permutation is applied across images for a run. The current saved results include the unpermuted case, represented as one tile, and 3-by-3 grids, represented as nine tiles. The original proposal also illustrates 4-tile and 9-tile permutations as target resolutions for this line of experimentation.

4 Methodology

The experimental pipeline is organized into three notebooks. Part 1 establishes the ResNet-18 baseline by training on the unpermuted task and on fixed 3-by-3 tile permutations. Part 2 reuses the Part 1 baseline and trains improvement ablations on the same task family. Part 3 computes permutation descriptors and correlates them with the validation accuracy observed for the corresponding runs.

The primary model in the current results is ResNet-18 [3]. Experiments save raw run outputs, aggregated CSV summaries, and figures under `outputs/results` and `outputs/figures`. Aggregation reports mean final validation accuracy, mean best validation accuracy, standard deviation where multiple runs exist, and the number of runs.

4.1 Training and Validation Sets by Experiment

All experiments use the labeled Kaggle Dogs vs Cats training folder as the only source

Part	1	Training set	Validation set
Baseline		20,000 unpermuted images.	5,000 unpermuted held-out images.
4-by-4 permutations		Same 20,000 images with one fixed 4-by-4 tile permutation per run; five permutation IDs are saved.	Same 5,000 held-out images with the matching 4-by-4 permutation.
10-by-10 permutations		Same 20,000 images with one fixed 10-by-10 tile permutation per run; five permutation IDs are saved.	Same 5,000 held-out images with the matching 10-by-10 permutation.

Table 1: Part 1 training and validation sets. The same split is reused for ResNet-18, DeiT-Tiny, and MLP-Mixer Base.

of supervised examples. Labels are inferred from filenames, and the samples are split once with a stratified class-balanced split using seed 42, validation fraction 0.2, and test fraction 0.0. For the full dataset this gives 20,000 training images and 5,000 validation images: 10,000 cats and 10,000 dogs for training, and 2,500 cats and 2,500 dogs for validation. The validation split is held out from optimization and is not modified by training-only augmentations such as random erasing, CutMix, patch shuffle, or curriculum-stage sampling.

The curriculum ablations in Part 2 use stage-specific training loaders, but they do not create separate held-out validation splits per stage. Across both curricula, validation remains the same 5,000-image held-out loader for the final target run.

TODO: Synchronize the experiment lists, tile sizes, tables, figures, and narrative claims

across the abstract, data overview, methodology, results, discussion, and conclusion so that every section describes the same current experiment matrix.

4.2 Architecture Comparison and Frozen Backbones

To compare architectural inductive biases beyond ResNet-18, the Part 1 pipeline also supports three pretrained feature extractors: ResNet-18, DeiT-Tiny [5], and MLP-Mixer [4]. All three are used in the same transfer-learning regime: the ImageNet-pretrained backbone is frozen and only the final binary classification head is optimized. This makes the relevant training-capacity comparison the number of learnable head parameters, not the total number of parameters stored in the network.

The three models differ mainly in how they exchange information across the image. ResNet-18 uses convolutions, giving it a strong local-neighborhood bias. DeiT-Tiny represents the image as patches and uses self-attention to let patches interact. MLP-Mixer also operates on image patches, but replaces attention and convolutions with two MLP operations: token mixing, which combines information across patches, and channel mixing, which recombines feature dimensions inside each patch representation.

The MLP-Mixer choice required special care. The smaller `mixer_s16_224` architecture would have been closer in total parameter count, but the available `timm` configuration does not provide pretrained weights for that variant. Freezing an untrained MLP-Mixer backbone would train only a classifier on random fixed features, which would not be a meaningful comparison to pretrained ResNet-18 and DeiT-Tiny. We therefore use the larger pretrained




Figure 1: Simplified MLP-Mixer flow. Token mixing exchanges information across patches, while channel mixing recombines features within each patch representation.

`mixer_b16_224.goog_in21k_ft_in1k` model. Although this increases the number of frozen parameters and may increase forward-pass cost, the number of optimized parameters remains comparable to the other models: 1,538 learnable parameters for MLP-Mixer Base versus 1,026 for ResNet-18 and 386 for DeiT-Tiny.

4.3 Improvement Strategies

Part 2 evaluates training strategies against the regular Part 1 ResNet-18 baseline. The main distinction is whether a strategy adds new training views of the same images, changes the optimization objective, or changes the order in which the model sees easy and hard inputs. In all cases, validation is performed on the held-out validation images with the target permutation and without training-only perturbations.

4.3.1 Regular Part 1 Baseline

Main idea. This is the reference point for Part 2. ResNet-18 is trained normally on the target permutation, with no extra augmentation, no curriculum, and the standard loss.

Reason. The baseline is necessary because every ablation must be judged against ordi-

nary training on the same task. Without it, an improvement strategy might look useful simply because the target permutation is easier than another run.

Deeper dive. The baseline uses the same train and validation split as Part 1. For a permuted run, both training and validation images are transformed by the same fixed tile permutation for that run. It has one training stage and does not increase the dataset size.

4.3.2 Patch Shuffle

Main idea. Patch shuffle randomly rearranges smaller patches during training. The model therefore sees extra local spatial disruptions in addition to the fixed tile permutation.

Reason. This is relevant because the main task already tests whether class information survives spatial reordering. If extra random patch movement helps, it suggests that robustness to local layout disruption is useful for permuted Dogs versus Cats classification.

Deeper dive. The target permutation is still the run-level task definition, but each training batch can also receive random patch shuffling. This does not duplicate image files or increase the stored dataset beyond 20,000 training images; instead, it creates different stochastic views during optimization. The validation loader does not use patch shuffle, so reported accuracy measures the fixed target validation permutation.

4.3.3 Random Erasing

Main idea. Random erasing hides random image regions during training. The model must learn from the remaining visible evidence rather than depending on one part of the animal.

Reason. Tile permutations can break global

shape, so a model may need to rely on partial cues such as fur texture, ears, eyes, color, or background context. Random erasing tests whether forcing the model to survive missing local evidence improves that behavior.

Deeper dive. This method creates occluded variants of the same training examples on the fly. It does not store a larger dataset; each epoch can expose different erased regions from the same images. The validation set remains un-erased and is transformed only by the target permutation.

4.3.4 Same-Label CutMix

Main idea. Same-label CutMix replaces part of one training image with a patch from another image of the same class. The label stays unchanged because cat is mixed only with cat and dog only with dog.

Reason. This experiment asks whether more same-class local variation helps the model classify permuted images. It is especially relevant when global object layout is damaged and the classifier may need to combine local evidence from separated tiles.

Deeper dive. Unlike ordinary CutMix, this version avoids mixed labels. The added training views are hybrid same-class images produced during batching, not a new stored dataset. Validation uses the original held-out images and does not apply CutMix.

4.3.5 Combined Augmentations

Main idea. Combined augmentations apply more than one training perturbation in the same ablation. The goal is to expose the model to richer variation than any single augmentation provides.

Reason. Different augmentations stress different weaknesses: patch shuffle stresses spatial order, CutMix stresses partial object ev-

idence, and image perturbations stress reliance on specific pixels. Combining them tests whether these robustness pressures are complementary.

Deeper dive. The training split remains the same 20,000 images, but the effective variety of examples increases because stochastic augmentations produce multiple possible views. The validation set is not augmented; it remains the fixed held-out target-permutation set. This makes the comparison a test of whether richer training exposure improves ordinary target validation accuracy.

4.3.6 Focal Loss

Main idea. Focal loss changes the training objective so harder examples receive more emphasis and easy examples receive less.

Reason. Some permuted images may remain easy because local class cues are obvious, while others may become much harder after tile movement. Focal loss tests whether concentrating learning on those harder cases improves performance.

Deeper dive. This strategy does not add new images to the training set and has one stage. It changes only the loss used during optimization. Validation uses the same held-out target set, so any difference comes from the training objective rather than from extra data.

4.3.7 Corruption-Probability Curriculum

Main idea. The model starts with mostly intact images and gradually sees the target permutation more often. The corruption probability increases across four stages.

Reason. Jumping immediately into a hard permutation may make optimization unstable. This curriculum tests whether gradual

exposure lets the model first learn dog-versus-cat cues and then adapt those cues to more frequent spatial corruption.

Deeper dive. Each stage uses the same 20,000 training images and target tile permutation, but changes the probability that the permutation is applied. The stages use probabilities 0.1, 0.3, 0.5, and 0.8. This does not store a larger dataset; it changes the sampled training view. There is no separate validation set for each stage; validation always uses the fixed target validation loader.

4.3.8 Permutation-Difficulty Curriculum

Main idea. The model starts on easier spatial structure and moves toward the final harder permutation. Instead of changing probability, it changes the tile grid difficulty across stages.

Reason. Smaller or easier permutations preserve more of the original image structure. This experiment tests whether a model can learn useful class cues under mild disruption before facing the final, harder tile rearrangement.

Deeper dive. Each stage uses the same 20,000 training images but changes the tile grid used for training. The schedule begins with original images and then progresses through 2-by-2, 3-by-3, and 4-by-4 permutations; for 10-by-10 target runs, it adds a final 10-by-10 stage. The stored dataset size does not increase, and validation remains the held-out loader for the final target run.

4.4 Tile Permutation Hardness Metrics

We quantify the difficulty induced by a tile permutation using three complementary, model-agnostic axes and a fourth combined

score. Let the image be divided into an $N \times N$ grid. For each source tile i , let (r_i, c_i) denote its original row and column, and let (r'_i, c'_i) denote its destination after applying the permutation.

4.4.1 Adjacency Destruction Hardness

Main idea. This metric asks how much of the original neighborhood structure survives after the tile permutation. If tiles that were next to each other are no longer neighbors, the image becomes harder for a convolutional model that relies on local structure.

For every original rightward and downward neighboring relation, we test whether the same two source tiles remain adjacent in the same direction after permutation. The number of such canonical adjacency relations is

$$M = 2N(N - 1).$$

If $A_{\text{preserved}}$ is the number of preserved directed adjacencies, adjacency destruction is

$$A = 1 - \frac{A_{\text{preserved}}}{M}.$$

A low value means that much of the original local structure is preserved. A high value means that neighboring tile relations were mostly destroyed.

4.4.2 Global Tile Displacement

Main idea. This metric measures how far tiles move from their original positions. A permutation that only nudges tiles nearby should preserve more global shape than one that sends tiles across the image.

For each tile, we compute its Manhattan displacement in grid coordinates:

$$\delta_i = |r'_i - r_i| + |c'_i - c_i|.$$

The normalized global displacement is

$$D = \frac{\sum_i \delta_i}{D_{\text{max}}},$$

where D_{max} is the maximum possible total Manhattan displacement for an $N \times N$ grid. A low value indicates that tiles remain close to their original locations. A high value indicates large-scale spatial relocation.

4.4.3 Spatial Permutation Entropy

Main idea. This metric asks whether tile movement follows a simple pattern or many different movement patterns. More varied movement directions and distances mean the permutation is less orderly.

For each tile, we compute

$$\Delta_i = (r'_i - r_i, c'_i - c_i)$$

and assign it to a movement bin

$$b_i = (\text{direction}(\Delta_i), \text{distanceTier}(\Delta_i)).$$

The direction is one of

$$\{\text{stationary}, N, S, E, W, NE, NW, SE, SW\}.$$

For non-stationary tiles, the normalized Manhattan distance is

$$d_i = \frac{|r'_i - r_i| + |c'_i - c_i|}{2(N - 1)}.$$

The distance tier is near when $0 < d_i \leq 1/3$, medium when $1/3 < d_i \leq 2/3$, and far when $2/3 < d_i \leq 1$. Stationary tiles use a single stationary bin. If p_b is the empirical probability of movement bin b , the normalized entropy is

$$E = \frac{-\sum_b p_b \log p_b}{\log(N^2)}.$$

A low value means that movement is structured or concentrated in a small number of movement patterns. A high value means that the permutation contains diverse movement directions and distances, making it more disorderly.

4.4.4 Combined Hardness Score

Main idea. The combined score summarizes the previous metrics into one number so permutations can be ranked by an overall notion of difficulty.

The final score is a convex combination of the three normalized components:

$$C = w_a A + w_d D + w_e E,$$

with default weights

$$w_a = 0.5, \quad w_e = 0.3, \quad w_d = 0.2.$$

A low combined score represents an easier permutation with preserved structure, small movement, and low movement disorder. A high combined score represents a harder permutation across these complementary axes.

5 Results

Part 1 shows that 3-by-3 tile permutations reduce mean final validation accuracy from 86.5% to 83.7%, but do not make the task infeasible. This supports the proposal’s hypothesis that the Dogs versus Cats image setting is a useful relaxation: the permutation disrupts spatial layout, yet enough visual class information remains for a standard convolutional model to learn above chance.

Part 2 identifies permutation-difficulty curriculum training as the best final 3-by-3 setting in the saved results, reaching 89.4% mean final validation accuracy and 93.3% mean best validation accuracy. Patch shuffle is the strongest non-curriculum ablation at 87.5%. The corruption-probability curriculum has the highest mean best validation accuracy, 97.1%, but its final accuracy is lower, suggesting either instability, overfitting after the best epoch, or sensitivity to the short training configuration.




Figure 2: Part 1 validation accuracy as a function of tile count.

Part 3 provides an initial ranking lens for permutations. Adjacency destruction, spatial entropy, and combined hardness all have negative Pearson correlations with validation accuracy in the current joined data, matching the intuition that destroying local structure and increasing disorder make classification harder. The sample size is small, so these values should be interpreted as exploratory rather than conclusive.

6 Discussion

The proposal asks whether a neural network can classify post-permutation images and whether performance depends on tile resolution and permutation choice. The current results answer the first question affirmatively for 3-by-3 tile permutations: the ResNet-18 baseline loses accuracy relative to unpermuted images, but the drop is modest in the saved run configuration. This suggests that the model may rely on local texture, color, object parts, and class-specific visual cues that survive tile reordering.

The ablation results indicate that training strategy matters. Patch shuffle and permutation-difficulty curriculum improve over the regular 3-by-3 baseline, which is con-




Figure 3: Part 2 comparison of augmentation and curriculum strategies.




Figure 4: Part 3 metric-versus-accuracy plots for the configured hardness metrics.

sistent with the idea that exposure to spatial disruption can make the model more robust to permuted inputs. The gap between best and final accuracy for the corruption-probability curriculum is an important warning: a method can discover a strong checkpoint while still finishing poorly, so best-checkpoint reporting and final-epoch reporting should both remain in the evaluation.

The hardness metrics address the proposal’s third objective: ranking permutations with the same tile count. The current evidence is directionally sensible but underpowered. With only three joined records, a single run can strongly affect the measured correlations. A stronger conclusion will require more tile resolutions, more random permutations per resolution, repeated seeds, and possibly separate analysis for each model family.

7 Conclusion and Future Work

This project implements a reproducible Dogs versus Cats testbed for studying fixed tile-wise image permutations. The work translates the proposal’s motivation into three concrete outputs: baseline classification under permutation, improvement ablations for the permuted setting, and permutation hardness metrics that can be joined to validation accuracy. The current results show that 3-by-3 permutations do not destroy classifiability, that curriculum and spatial-disruption augmentation can improve performance, and that model-agnostic hardness metrics provide a promising way to rank permutations.

Future work should expand the experiment matrix beyond the current saved runs. The next steps are to evaluate additional tile resolutions such as 2-by-2 and larger grids, sample more permutations per resolution, repeat experiments across seeds, compare more ar-

chitectures, and test whether a model trained on one permutation transfers to another. These extensions would better separate memorization of a fixed permutation from direct use of class cues in permuted images, which is one of the central questions posed by the original project proposal.

Part 2 strategy	Training set	Added training exposure	Validation set
Regular Part 1 baseline	Reused Part 1 ResNet-18 runs.	None beyond the fixed target permutation.	Reused Part 1 held-out validation accuracy.
Patch shuffle	Same 20,000 training images and target permutation.	Extra batch-level random patch shuffling during training, so the model sees additional local rearrangements beyond the fixed dataset examples.	Same 5,000 held-out images with the target permutation only.
Random erasing	Same 20,000 training images and target permutation.	Randomly removes image regions during training, creating occluded variants of the original training images.	Same 5,000 held-out images with no random erasing.
Same-label Cut-Mix	Same 20,000 training images and target permutation.	Mixes image patches only between examples with the same label, adding same-class hybrid training images without changing labels.	Same 5,000 held-out images with no CutMix.
Combined augmentations	Same 20,000 training images and target permutation.	Combines multiple augmentation sources, such as patch shuffle and same-label CutMix, to create richer training variants.	Same 5,000 held-out images with no training-only augmentations.
Focal loss	Same 20,000 training images and target permutation.	No new images; the loss reweights hard examples during optimization.	Same 5,000 held-out images with the target permutation.
Corruption-probability curriculum	Same 20,000 training images at every stage.	Stages use target permutation probabilities 0.1, 0.3, 0.5, and 0.8.	One fixed target validation loader; no separate validation set per stage.
Permutation-difficulty curriculum	Same 20,000 training images at every stage.	Stages progress through original, 2-by-2, 3-by-3, 4-by-4, and for 10-by-10 targets an additional 10-by-10 stage.	One fixed final-target validation loader; no separate validation set per stage.

Table 2: Part 2 training and validation sets by ablation. The current output matrix uses ResNet-18 with the baseline, three 4-by-4 permutations, and three 10-by-10 permutations where available.

Part 2 strategy	Stages	Experiment sequence	Dataset increase
Regular Part 1 baseline	1	Standard training on the target permutation.	No increase: 20,000 training examples per epoch.
Patch shuffle	1	Standard training, with random patch shuffling applied inside training batches.	No stored dataset growth; stochastic batch augmentation creates new views on the fly.
Random erasing	1	Standard training, with random image-region erasing in the training transform.	No stored dataset growth; each epoch can show different occluded variants.
Same-label Cut-Mix	1	Standard training, with same-class patches mixed inside training batches.	No stored dataset growth; batch mixing creates extra hybrid views during optimization.
Combined augmentations	1	Standard training, with more than one augmentation source active together.	No stored dataset growth; effective exposure increases through multiple stochastic views.
Focal loss	1	Standard training with focal loss instead of cross-entropy.	No increase: the examples are unchanged, only their loss weights change.
Corruption-probability curriculum	4	Permutation probability stages: 0.1, 0.3, 0.5, 0.8.	No stored dataset growth; each stage reuses all 20,000 images with a different corruption probability.
Permutation-difficulty curriculum	1, 4, or 5	Baseline target: original only. 4-by-4 target: original, 2-by-2, 3-by-3, 4-by-4. 10-by-10 target: original, 2-by-2, 3-by-3, 4-by-4, 10-by-10.	No stored dataset growth; each stage reuses all 20,000 images with a different difficulty level.

Table 3: Part 2 ablation experiment design. Augmentations do not duplicate files on disk; they increase the diversity of training views produced from the same fixed split.

Part 3 analysis	Training set	Validation source
Hardness metric computation	Not applicable; Part 3 does not train a model.	Reuses Part 1 ResNet-18 validation accuracies.
Metric-accuracy join	Not applicable; only saved CSV outputs are joined.	Matches each metric row to the corresponding Part 1 permutation ID.
Correlation analysis	Not applicable; no parameters are updated.	Computes correlations against the joined held-out validation accuracies.

Table 4: Part 3 data usage. Part 3 analyzes saved permutations and validation accuracies rather than creating new train or validation splits.

Model	Main operation	Inductive bias
ResNet-18	Convolution	Local texture and translation structure
DeiT-Tiny	Self-attention	Patch-to-patch interactions
MLP-Mixer Base	MLP mixing	Patch and feature mixing without attention

Table 5: Architectural comparison of the three frozen feature extractors.

Model	Learnable	Frozen	Total
DeiT-Tiny	386	5,524,416	5,524,802
ResNet-18	1,026	11,176,512	11,177,538
MLP-Mixer Base	1,538	59,111,472	59,113,010

Table 6: Parameter counts for the frozen-backbone comparison with two output classes. “Learnable” counts only the classifier head optimized during training; “Frozen” counts pretrained parameters used only in the forward pass.

Metric	Dimension	What it captures
Adjacency destruction hardness	Local topology	Are original neighbors still neighbors?
Global tile displacement	Global geometry	How far did tiles move from their original position?
Spatial permutation entropy	Permutation disorder	How diverse or irregular are tile movement patterns?
Combined hardness score	Aggregated hardness	Weighted summary of the three normalized axes

Table 7: Tile permutation hardness metrics used in the analysis.

Setting	Tiles	Final Acc.	Runs
Unpermuted	1	86.5%	1
3-by-3 permuted	9	83.7%	2

Table 8: Part 1 ResNet-18 validation accuracy aggregated from saved experiment outputs.

Ablation	Tiles	Final Acc.	Best Acc.	Runs
Regular Part 1 baseline	9	83.7%	83.7%	2
Combined augmentations	9	81.7%	81.7%	2
Patch shuffle	9	87.5%	87.5%	2
Random erasing	9	83.7%	83.7%	2
Same-label CutMix	9	84.6%	84.6%	2
Corruption-probability curriculum	9	78.8%	97.1%	2
Permutation-difficulty curriculum	9	89.4%	93.3%	2

Table 9: Part 2 3-by-3 ablation results for ResNet-18.

Hardness metric	Pearson	Spearman
Global displacement	-0.04	0.50
Adjacency destruction	-0.47	-0.50
Spatial entropy	-0.39	-0.50
Combined hardness	-0.38	-0.50

Table 10: Part 3 correlations between ResNet-18 validation accuracy and permutation metrics, computed over three joined records.

References

- [1] E. Boyle, J. Holmgren, and M. Weiss. Permuted puzzles and cryptographic hardness. In *Theory of Cryptography*, pages 465–493, 2019.
- [2] I. Goodfellow, Y. Bengio, and A. Courville. *Deep Learning*. MIT Press, 2016.
- [3] K. He, X. Zhang, S. Ren, and J. Sun. Deep residual learning for image recognition. In *IEEE Conference on Computer Vision and Pattern Recognition*, pages 770–778, 2016.
- [4] I. O. Tolstikhin, N. Houlsby, A. Kolesnikov, L. Beyer, X. Zhai, T. Unterthiner, J. Yung, A. Steiner, D. Keysers, J. Uszkoreit, M. Lucic, and A. Dosovitskiy. MLP-Mixer: An all-MLP architecture for vision. In *Advances in Neural Information Processing Systems*, 2021.
- [5] H. Touvron, M. Cord, M. Douze, F. Massa, A. Sablayrolles, and H. Jégou. Training data-efficient image transformers and distillation through attention. In *International Conference on Machine Learning*, pages 10347–10357, 2021.